

Faizan Tanveer

📞 +92-3252726321 | ✉ faizantanveerai@gmail.com | 🌐 faizantanveerr | 🏠 faizantanveeer | 📍 Lahore, Pakistan

Summary

AI/ML Engineer focused on building practical and scalable Machine Learning and Generative AI solutions using LLMs. Experienced in designing end to end ML pipelines, developing reliable RAG systems, and improving retrieval quality to deliver more accurate and context-aware AI applications. Strong foundation in fine tuning, embeddings, prompt engineering, and deploying production grade solutions on Microsoft Azure and Google Cloud Platform, particularly within Legal and Health Tech domains. Previously contributed at **Tkxel**, where I strengthened data workflows and automated reporting processes, enabling teams to make more reliable and timely decisions. I bring a problem solving mindset and a strong sense of ownership, aiming to create AI systems that deliver measurable business value.

Experience

Associate AI/ML Engineer

Dec 2025 – Present

Hazen Technologies Inc.

Lahore, Pakistan

- Used **Azure Document Intelligence** to extract structured patient data from medical ledgers for legal case intake at Ged Lawyers, fully replacing manual data entry for a law firm client.
- Engineered **Syntellio**, a multi-agent RAG system with MCP support, achieving **60%+ retrieval accuracy** and **35% lower query latency** via semantic search, metadata filtering, and similarity scoring.
- Revamped legacy codebase to **agentic AI architecture** using LangChain agents and MCP tooling, enabling autonomous multi-step reasoning across Legal and Health Tech workflows.
- Designed **client-facing analytics dashboards** integrated with live production pipelines for real-time operational reporting.

Business Intelligence Analyst

Aug 2025 – Dec 2025

Tkxel

Lahore, Pakistan

- Eliminated **35+ hours/week** of manual reporting by consolidating 5+ dashboards into a unified real-time analytics suite (Power BI, Looker, DAX).
- Designed **production-grade ETL pipelines** on BigQuery processing **90GB+ multi-source data**; boosted KPI accuracy by **25%** via advanced data cleaning and feature engineering.
- Accelerated reporting speed by **30%** with performance-tuned dashboards serving marketing, finance, and operations teams.
- Automated ETL pipelines for preprocessing data with Python & SQL, cutting data prep time by **40%** and enabling faster ML-ready dataset delivery.

Projects

Syntellio – Multi Agent RAG System | Python, LangChain, MCP, LLMs, FAISS, Pinecone

Hazen Technologies

- Addressed inconsistent retrieval quality and shallow single step reasoning in existing document Q&A systems by architecting a multi agent RAG framework with structured multi step decision flows.
- Redesigned static pipelines into a scalable agent driven architecture, improving retrieval accuracy by **60%+** and reducing response latency by **35%**.
- Solved context fragmentation issues in Legal and Health Tech documents by improving grounding and answer consistency across large knowledge bases.

InstaSum – Instagram Reels AI Summarizer | Flask, Whisper, GPT, Azure OpenAI, SQLite, SSE

Personal Project

- Solved the challenge of extracting structured insights from short form video content by building an AI pipeline that automates transcription and summarization.
- Mitigated frequent API failures and rate limits by designing a three layer resilient fetching strategy, reducing failure rates by **40%+** and ensuring consistent content retrieval.
- Enabled real time processing and scalable usage through modular API design and streaming architecture.

Technical Skills

Python, SQL, Machine Learning, Deep Learning, NLP, LLMs, Generative AI, RAG, Agentic AI, Multi-Agent Systems, MCP, LangChain, LangGraph, Azure, GCP, AWS, ETL Pipelines, Power BI, Looker Studio, Git, Docker

Education

University of Management and Technology (UMT)

Lahore, Pakistan

Bachelor of Science in Computer Science (BSCS) | **CGPA: 3.66 / 4.0**

Oct 2021 – Jun 2025