

# Faizan Tanveer

📞 +92-3252726321 | ✉ faizantanveerai@gmail.com | 🌐 faizantanveerr | 🏠 faizantanveeer | 📍 Lahore, Pakistan

## Summary

AI/ML Engineer focused on building scalable Machine Learning and Generative AI solutions using Large Language Models (LLMs). Experienced in designing end-to-end ML pipelines, Retrieval-Augmented Generation (RAG) systems, fine-tuning, embeddings, and prompt engineering. Proficient in deploying production-grade solutions on Microsoft Azure and Google Cloud Platform (GCP). Passionate about building intelligent, context-aware AI applications that solve real-world problems across diverse domains. Previously at Txxel, automated reporting and strengthened data workflows enabling reliable and timely decision-making. Ownership-driven, focused on delivering measurable business value.

## Experience

### Associate AI/ML Engineer

Dec 2025 – Present

Hazen Technologies Inc.

Lahore, Pakistan

- Used **Azure Document Intelligence** to extract structured patient data from medical ledgers for legal case intake at Ged Lawyers, fully replacing manual data entry for a law firm client.
- Engineered **Syntellio**, a multi-agent RAG system with MCP support, achieving **60%+ retrieval accuracy** and **35% lower query latency** via semantic search, metadata filtering, and similarity scoring.
- Revamped legacy codebase to **agentic AI architecture** using LangChain agents and MCP tooling, enabling autonomous multi-step reasoning across Legal and Health Tech workflows.
- Designed **client-facing analytics dashboards** integrated with live production pipelines for real-time operational reporting.

### Business Intelligence Engineer

Aug 2025 – Dec 2025

Txxel

Lahore, Pakistan

- Eliminated **35+ hours/week** of manual reporting by consolidating 5+ dashboards into a unified real-time analytics suite (Power BI, Looker, DAX).
- Designed **production-grade ETL pipelines** on BigQuery processing **90GB+ multi-source data**; boosted KPI accuracy by **25%** via advanced data cleaning and feature engineering.
- Accelerated reporting speed by **30%** with performance-tuned dashboards serving marketing, finance, and operations teams.
- Automated ETL pipelines for preprocessing data with Python & SQL, cutting data prep time by **40%** and enabling faster ML-ready dataset delivery.

## Projects

### Syntellio – Multi Agent RAG System | Python, LangChain, MCP, LLMs, FAISS, Pinecone

Hazen Technologies

- Addressed inconsistent retrieval quality and shallow single-step reasoning in existing document Q&A systems by architecting a multi-agent RAG framework with structured multi-step decision flows.
- Redesigned static pipelines into a scalable agent-driven architecture, improving retrieval accuracy by **60%+** and reducing response latency by **35%**.
- Solved context fragmentation issues in Legal and Health Tech documents by improving grounding and answer consistency across large knowledge bases.

### InstaSum – Instagram Reels AI Summarizer | Flask, Whisper, Azure OpenAI, Supabase

Personal Project

- Solved the challenge of extracting structured insights from short-form video content by building an AI pipeline that automates transcription and summarization.
- Mitigated frequent API failures and rate limits by designing a three-layer resilient fetching strategy, reducing failure rates by **40%+** and ensuring consistent content retrieval.
- Enabled real-time processing and scalable usage through modular API design and streaming architecture.

## Technical Skills

Python, SQL, NoSQL, Machine Learning, Deep Learning, NLP, Large Language Models (LLMs), Generative AI, Retrieval-Augmented Generation (RAG), Fine-Tuning, Prompt Engineering, Agentic AI, Multi-Agent Systems, LangChain, LangGraph, MCP, ETL Pipelines, Azure, GCP, AWS, Power BI, Looker Studio, Git

## Education

### University of Management and Technology (UMT)

Bachelor of Science in Computer Science (BSCS) | CGPA: 3.66 / 4.0

Lahore, Pakistan

Oct 2021 – Jun 2025